



Majhi OS
majhi.tech

RESEARCH REPORT · 2025

The Operational Visibility Gap

Why recruiting breaks at scale — and what the infrastructure layer that fixes it actually looks like. A research report for senior talent and hiring leaders.

68%

VP searches stall past week 10 due to execution failure

14wk

Avg time-to-fill on fragmented tooling

\$3.8K

Monthly tool cost per recruiter before first hire

50d

Avg close time on unified infrastructure (Majhi OS)



EXECUTIVE SUMMARY

The problem isn't talent scarcity.

01

Every senior recruiting leader has experienced it: a search that should close in eight weeks stretches to twenty. A mandate progressing well suddenly has no candidates in late-stage evaluation. A shortlist that passed internal review gets rejected by the board for lack of evidence.

The primary cause is operational invisibility. Recruiting teams are running critical processes through fragmented stacks of tools that share no state, produce no unified signal, and offer no systemic view of what's happening inside a search. This report documents the structural failure modes, proposes an operational maturity model, and describes what genuine hiring infrastructure actually requires.




Stall Rate

VP searches stalling past week 10 on fragmented tooling



TtF Reduction

Faster time-to-fill vs. fragmented average (VP level)



Audit Trail Coverage

Fragmented teams with full decision audit trails



The difference between a search that closes in 50 days and one that closes in 120 is rarely the talent market. It's almost always the execution layer — and whether anyone can see inside it.

— Manas Majhi, Founder, Majhi OS




SECTION 01 · BENCHMARK DATA

The true cost of fragmented hiring


02

Time-to-Fill: VP-Level Role Comparison



Majhi OS clients close VP-level mandates in an average of 50 days vs. the 14-week fragmented average. Based on 2024 engagements.

Monthly Tooling Cost per Recruiter (Before First Hire)



Majhi OS unified pricing vs. typical SMB vendor pricing at standard seat counts. Savings vary by volume.



SECTION 02 · ROOT CAUSE ANALYSIS

Four structural failure modes of fragmented hiring

03

These failure modes appear predictably regardless of team size, seniority, or role complexity. They are not symptoms of bad recruiting — they are symptoms of missing infrastructure.

01

Intake without normalization

Candidates arrive from LinkedIn exports, CSV uploads, referrals, and ATS imports in different formats with no deduplication. The same candidate may enter the funnel twice with no system to detect or resolve the collision. Quality gates can't be applied to unstandardized profiles — every hour spent on a duplicate is an hour not spent on execution.

02

Outreach without verification

Email sequences are built on unverified contact data. Bounces damage domain reputation. Stale LinkedIn profiles generate outreach to candidates who left their roles months ago. When sequences run without DNS/MX verification and bounce-kill logic, the entire outreach operation runs on a slowly degrading foundation — and no one knows until reply rates collapse.

03

Evaluation without evidence

Hiring recommendations are made on résumés, informal conversations, and recruiter judgment — none of which is systematically recorded or auditable. When a board asks "why this candidate?", the honest answer is often "because they interviewed well." Without structured evidence — proof points, risk flags, mandate alignment scores — every recommendation is an opinion dressed as analysis.

04

Execution without observability

The most dangerous failure mode. A search can be technically in progress while the pipeline is functionally stalled. Without real-time mandate health states, funnel velocity metrics, and stall detection, no one knows a search is in trouble until it has already lost two to four weeks. By the time leadership asks for a status update, the data is spread across six tools — none of which talk to each other.

▲ THE COMPOUNDING EFFECT

These four failure modes do not operate independently. Intake failure produces evaluation failure. Evaluation failure corrupts outreach. Execution failure prevents observability. Each failure mode amplifies the others — which is why searches that "should" close in eight weeks regularly take twenty.



SECTION 03 · MATURITY MODEL

Fragmented operation vs. infrastructure operation

04

The gap between these two operating models is not primarily technological. Most recruiting teams have access to good individual tools. The gap is **architectural** — whether those tools share state, produce a unified signal, and enable systemic visibility into what's happening inside a search.


DIMENSION	✗ FRAGMENTED OPERATION	✓ INFRASTRUCTURE (MAJHI OS)
Intake	Multiple formats, no normalization, manual deduplication required	Normalized intake with duplicate detection and profile stitching across sources
Evaluation	Résumé + interview notes, no structured evidence or mandate scoring	Evidence-backed dossiers with proof points, risk flags, and mandate fit scores
Outreach	Unverified contacts, manual follow-up, no bounce logic or reply detection	DNS/MX verified sequences with stop-on-reply and automatic bounce-kill
Pipeline health	Spreadsheet + weekly status calls; problems surface 2–3 weeks late	Live mandate health states with funnel decay alerts and stall detection
Reporting	Manual exports, always retrospective, no real-time mandate signal	Real-time telemetry across all active mandates in one command surface
Auditability	No record of recruiter actions; decisions untraceable if challenged	Full audit log — every event, action, and override timestamped and searchable

Key Performance Indicators: Fragmented vs. Majhi OS

Fragmented Stack Majhi OS


Time-to-Fill (VP Level)

Median weeks across mandates



Outreach Reply Rate

With vs. without DNS/MX verification




Shortlist Approval Rate

Board / CEO accepted shortlists



Searches with Full Audit Trail

Timestamped decision record end-to-end





SECTION 04 · PIPELINE ANALYSIS

Where candidates are lost in the funnel

05

The funnel comparison below illustrates how fragmented tooling creates compounding attrition at every stage. Majhi OS's normalized intake, QMS gating, and verified outreach retain significantly more qualified candidates from sourcing through to hire.

Recruitment Funnel: Fragmented Stack vs. Majhi OS (per 400 sourced candidates)

STAGE	FRAGMENTED STACK	MAJHI OS	
Sourced	400	400	
Normalized / QMS	210	380	+81%
Dossier Ready	160	365	+128%
Outreach Sent	140	340	+143%
Response Received	48	160	+233%
Interview Stage	24	90	+275%
Offer Made	12	60	+400%
Hire	8	52	+550%

Illustrative funnel. Fragmented reflects unmanaged drop-off; Majhi OS reflects gated, evidence-backed execution. Based on median VP-level search data from 40 mandates.

WHY THE FUNNEL DIVERGES SO SHARPLY

The compounding effect of four failure modes becomes visible in funnel analysis. At every stage, fragmented tooling introduces attrition that Majhi OS prevents: normalized intake retains 95% vs. 52% of sourced candidates into the QMS; verified outreach delivers 3× the reply rate; evidence-backed shortlists produce 2× the approval rate. By hire, the efficiency gap is 6×.



SECTION 05 · ARCHITECTURE

What infrastructure-grade execution looks like

06

Infrastructure is not a faster ATS. It is a system that owns the execution loop end-to-end — with gated handoffs between stages, evidence anchoring at every decision point, and real-time observability across the entire operation. The recruiter doesn't manage the process. They govern it.

```
majhi-os · execution-loop · mandate-0041 · vp-engineering

01  Candidate PDF uploaded · intake normalized · format standardized · duplicate check passed
02  Identity stitched across LinkedIn + CRM + prior searches · no collision detected
03  QMS gate · fit score 8.7/10 · mandate alignment: STRONG · advanced to dossier
04  Dossier auto-generated · 7 proof points extracted · 2 risk flags surfaced for review
05  [ HUMAN ] Recruiter review · dossier approved · outreach sequence unlocked
06  3-touch sequence queued · DNS/MX verified · bounce-kill active · stop-on-reply enabled
07  [ ALERT ] Response rate -28% vs. mandate baseline · stall probability: 74% · escalation queued
08  [ HUMAN ] Manual override · sequence paused · reason logged · alternative approach initiated
09  [ ALERT ] Pipeline telemetry updated · mandate health: AT RISK · leadership dashboard notified
```

DESIGN PRINCIPLES

PRINCIPLE 01

Evidence before opinion.

Every recommendation must be traceable to structured evidence — proof points, fit scores, risk flags. Unrecorded judgment is a liability at board level.

PRINCIPLE 02

Observability is foundational.

If you can't see what's happening inside a search in real time, you can't manage it. The command center is the primary control surface — not a reporting tool.

PRINCIPLE 03

Gate before you execute.

Every candidate passes through a quality gate before outreach. The QMS is what makes scale trustworthy, not just fast.

PRINCIPLE 04

One system. No handoffs.

Every tool handoff loses context and slows execution. Infrastructure means one connected system where data flows unbroken from intake to close.

PRINCIPLE 05

Automation reduces load.

The test for any automation: does it give the recruiter more clarity, or more to manage? Complexity without clarity is a more expensive version of the problem.

PRINCIPLE 06

Surface failures first.

Real infrastructure surfaces problems before humans discover them. Stall detection and funnel decay alerts are the primary signal the system is working.

CONCLUSION

Infrastructure thinking, not tool thinking.

07

The recruiting industry has invested heavily in tools that make individual tasks faster — better search, better enrichment, better sequencing. These are genuinely useful. But they do not solve the operational visibility gap, because the gap is not caused by slow individual tasks.

The gap is caused by the absence of a connected system that can see across all tasks simultaneously. Closing it requires a different architectural decision: to treat recruiting execution as an operational system, not a collection of workflows. To build around observability, not just efficiency. To make every candidate event, recruiter action, and quality gate decision visible, logged, and auditable.

Controlled execution is the only kind that scales without breaking. That is the commitment Majhi OS was built to keep.

50_d

Average VP close time on
Majhi OS

↓ 50% vs. avg

35%

Outreach reply rate with
verification

↑ 2.5× baseline

82%

Shortlist approval rate,
board/CEO

↑ from 38%

100%


Mandates with full audit trail

Every decision logged

SEE THE INFRASTRUCTURE IN ACTION

Walk through the command center, event stream, dossier system, and mandate health layer on a live brief with your context. Mission Walkthroughs are 45 minutes, with a senior Majhi OS team member.

Book at majhi.tech · support@majhi.tech



Methodology note: All benchmark data is derived from Majhi OS internal mandate data (2024–25) and publicly available industry research on executive recruiting timelines. VP-level is defined as VP, SVP, EVP, and C-suite roles. Funnel data is illustrative based on median searches. Individual results vary based on role, market, and team configuration.